

CONGRATULATIONS ISHAAN WATTS!

Ishaan did an amazing job in ramping up on the Wiz inference code. Thanks to him we were able to verify correctness in Gemax and setup a performant Wiz server for CALM. He also tried various quantization strategies to improve latency.

Furthermore, his latency benchmarking has been invaluable. It has given us clear, actionable insights for future optimizations. His contributions are deeply appreciated!!

